

Tiled Matrix Multiplication Grade

Grade Summary (History) (/grade/history/5983)

Created:	less than a minute ago (2022-03-05 09:05:25 +0000 UTC)
Total Score:	100 out of 100 points
Coding Score:	100 out of 100 points
Questions Score:	0

Program Code

```
1  #include <wb.h>
2
3  #define wbCheck\(stmt\)
4      do {
5          cudaError\_t err = stmt;
6          if \(err != cudaSuccess\) {
7              wbLog\(ERROR, "Failed to run stmt ", #stmt\);
8              wbLog\(ERROR, "Got CUDA error ... ", cudaGetErrorString\(err\)\);
9              return -1;
10         }
11     } while (0)
12
13  #define TILE\_WIDTH 32
14  #define BLOCK\_SIZE 8
15  int ceil\(int a, int b\){
16      return \(a + b - 1\)/b;
17  }
18
19
20  // Compute C = A * B
21  \_\_global\_\_ void matrixMultiplyShared\(float \*A, float \*B, float \*C,
```

```
22         int numARows, int numAColumns,
23         int numBRows, int numBColumns,
24         int numCRows, int numCColumns) {
25     //@@ Insert code to implement matrix multiplication here
26     //@@ You have to use shared memory for this MP
27     __shared__ float TileP[TILE_WIDTH][TILE_WIDTH];
28     __shared__ float TileQ[TILE_WIDTH][TILE_WIDTH];
29
30     int block_x = blockIdx.x;
31     int block_y = blockIdx.y;
32     int thread_x = threadIdx.x;
33     int thread_y = threadIdx.y;
34
35     int row = (block_y * blockDim.y) + thread_y;
36     int col = (block_x * blockDim.x) + thread_x;
37     int next = (numAColumns + BLOCK_SIZE - 1) / BLOCK_SIZE;
38     float hold = 0;
39
40     for (int a = 0; a < next; a++){
41         //load the first matrix tile
42         if((a * BLOCK_SIZE + thread_x) > numAColumns){
43             TileP[thread_y][thread_x] = 0.0;
44         }
45         else{
46             TileP[thread_y][thread_x] = A[row * numAColumns + a * BLOCK_SIZE +
47         }
48         //load the second matrix tile
49         if((a * BLOCK_SIZE + thread_y) >= numBRows){
50             TileQ[thread_y][thread_x] = 0.0;
51         }
52         else{
53             TileQ[thread_y][thread_x] = B[(a * BLOCK_SIZE + thread_y) * numBCol
54         }
55
56         __syncthreads();
57         //perform multiplication calculation
58         for (int b = 0; b < BLOCK_SIZE; b++){
59             hold += TileP[thread_y][b] * TileQ[b][thread_x];
60         }
61         __syncthreads();
62     }
63
64     if (row < numCRows && col < numCColumns){
65         C[row * numCColumns + col] = hold;
66     }
```

```
67
68     //__syncthreads();
69
70
71 }
72
73 int main(int argc, char **argv) {
74     wbArg_t args;
75     float *hostA; // The A matrix
76     float *hostB; // The B matrix
77     float *hostC; // The output C matrix
78     float *deviceA;
79     float *deviceB;
80     float *deviceC;
81     int numARows;    // number of rows in the matrix A
82     int numAColumns; // number of columns in the matrix A
83     int numBRows;    // number of rows in the matrix B
84     int numBColumns; // number of columns in the matrix B
85     int numCRows;    // number of rows in the matrix C (you have to set thi
86     int numCColumns; // number of columns in the matrix C (you have to set
87                     // this)
88
89     args = wbArg_read(argc, argv);
90
91     wbTime_start(Generic, "Importing data and creating memory on host");
92     hostA = (float *)wbImport(wbArg_getInputFile(args, 0), &numARows,
93                             &numAColumns);
94     hostB = (float *)wbImport(wbArg_getInputFile(args, 1), &numBRows,
95                             &numBColumns);
96     //@@ Set numCRows and numCColumns
97     numCRows = numARows;
98     numCColumns = numBColumns;
99     //@@ Allocate the hostC matrix
100    wbTime_stop(Generic, "Importing data and creating memory on host");
101
102    hostC = (float *) malloc((numCRows * numCColumns) * sizeof(float));
103
104    wbLog	TRACE, "The dimensions of A are ", numARows, " x ", numAColumns);
105    wbLog	TRACE, "The dimensions of B are ", numBRows, " x ", numBColumns);
106    wbLog	TRACE, "The dimensions of C are ", numCRows, " x ", numCColumns);
107
108    wbTime_start(GPU, "Allocating GPU memory.");
109    //@@ Allocate GPU memory here
110
111    int size_of_A = numARows * numAColumns * sizeof(float);
```

```
112 int size_of_B = numRows * numColumns * sizeof(float);
113 int size_of_C = numRows * numColumns * sizeof(float);
114 cudaMalloc((void **) &deviceA, size_of_A);
115 cudaMalloc((void **) &deviceB, size_of_B);
116 cudaMalloc((void **) &deviceC, size_of_C);
117
118 wbTime_stop(GPU, "Allocating GPU memory.");
119
120 wbTime_start(GPU, "Copying input memory to the GPU.");
121 //@@ Copy memory to the GPU here
122
123 cudaMemcpy(deviceA, hostA, size_of_A, cudaMemcpyHostToDevice);
124 cudaMemcpy(deviceB, hostB, size_of_B, cudaMemcpyHostToDevice);
125
126 wbTime_stop(GPU, "Copying input memory to the GPU.");
127
128 //@@ Initialize the grid and block dimensions here
129
130 dim3 dimensionBlock(BLOCK_SIZE, BLOCK_SIZE, 1);
131 dim3 dimensionGrid(ceil(numColumns, BLOCK_SIZE), ceil(numRows, BLOCK_
132
133 wbTime_start(Compute, "Performing CUDA computation");
134 //@@ Launch the GPU Kernel here
135
136 matrixMultiplyShared<<<dimensionGrid, dimensionBlock>>>(deviceA, device
137
138 cudaDeviceSynchronize();
139 wbTime_stop(Compute, "Performing CUDA computation");
140
141 wbTime_start(Copy, "Copying output memory to the CPU");
142 //@@ Copy the GPU memory back to the CPU here
143
144 cudaMemcpy(hostC, deviceC, size_of_C, cudaMemcpyDeviceToHost);
145
146 wbTime_stop(Copy, "Copying output memory to the CPU");
147
148 wbTime_start(GPU, "Freeing GPU Memory");
149 //@@ Free the GPU memory here
150
151 cudaFree(deviceA);
152 cudaFree(deviceB);
153 cudaFree(deviceC);
154
155 wbTime_stop(GPU, "Freeing GPU Memory");
156
```

```
157     wbSolution(args, hostC, numCRows, numCColumns);
158
159     free(hostA);
160     free(hostB);
161     free(hostC);
162
163     return 0;
164 }
165
```

Designed and architected by Abdul Dakkak (<https://www.dakkak.dev/>).