# **◆** (/mp/5984?show=code) List Reduction Attempt

## Attempt Summary

#### **Submit Attempt for Grading**

#### Remember to answer the questions before clicking.

Dataset Id: 5

Created: less than a minute ago (2022-04-11 03:17:43 +0000 UTC)

Status: Correct solution for this dataset.

## Timer Output

Kind	Location	Time (ms)	Message
Generic	main.cu::59	56.875034	Importing data and creating memory on host
GPU	main.cu::78	1.081262	Allocating GPU memory.
GPU	main.cu::85	0.053496	Copying input memory to the GPU.
Compute	main.cu::94	0.039577	Performing CUDA computation
Сору	main.cu::101	0.028169	Copying output memory to the CPU
GPU	main.cu::117	0.150213	Freeing GPU Memory

### Logger Output

webgpu.com/attempt/176392

2627

28

29

30 31

32 33

34

else

Level	Location	Message
Trace	main::75	The number of input elements in the input is 12670
Trace	main::76	The number of output elements in the input is 13

```
Program Code
    // MP Reduction
 1
    // Given a list (lst) of length n
    // Output its sum = lst[0] + lst[1] + ... + lst[n-1];
 4
   #include <wb.h>
 5
 6
 7
    #define BLOCK_SIZE 512 //@@ You can change this
 9
    #define wbCheck(stmt)
10
      do {
        cudaError_t err = stmt;
11
12
        if (err != cudaSuccess) {
          wbLog(ERROR, "Failed to run stmt ", #stmt);
13
          wbLog(ERROR, "Got CUDA error ... ", cudaGetErrorString(err));
14
15
          return -1;
16
17
      } while (0)
18
19
    __global__ void reduction(float *input, float *output, int len) {
20
      //@@ Load a segment of the input vector into shared memory
21
      __shared__ float sum[BLOCK_SIZE * 2];
22
23
      int thread_x = threadIdx.x;
24
25
      int load_in = 2*(BLOCK_SIZE * blockIdx.x) + thread_x;
```

webgpu.com/attempt/176392 2/5

for (int travel\_tree = BLOCK\_SIZE; travel\_tree >= 1; travel\_tree >>= 1)

sum[BLOCK\_SIZE + thread\_x] = input[load\_in + BLOCK\_SIZE];

sum[thread\_x] = load\_in < len ? input[load\_in] : 0.0;</pre>

if (load\_in + BLOCK\_SIZE < len)</pre>

//@@ Traverse the reduction tree

 $sum[BLOCK_SIZE + thread_x] = 0.0;$ 

```
__syncthreads();
35
          if (thread_x < travel_tree)</pre>
36
              sum[thread_x] += sum[thread_x + travel_tree];
37
38
     }
39
40
      //@@ Write the computed sum of the block to the output vector at the
41
      //@@ correct index
      if (thread_x == 0){
42
43
        output[blockIdx.x] = sum[0];
44
     }
45
   }
46
47
    int main(int argc, char **argv) {
48
      int ii;
49
     wbArg_t args;
50
     float *hostInput; // The input 1D list
     float *hostOutput; // The output list
51
52
     float *deviceInput;
     float *deviceOutput;
53
54
      int numInputElements; // number of elements in the input list
      int numOutputElements; // number of elements in the output list
55
56
57
      args = wbArg_read(argc, argv);
58
59
      wbTime_start(Generic, "Importing data and creating memory on host");
60
     hostInput = (float *)wbImport(wbArg_getInputFile(args, 0), &numInputEle
61
62
     numOutputElements = numInputElements / (BLOCK_SIZE << 1);</pre>
      if (numInputElements % (BLOCK_SIZE << 1)){</pre>
63
64
        numOutputElements++;
65
      }
66
67
      //2 variables to help with sizes
68
      int Size_of_Input = numInputElements * sizeof(float);
69
      int Size_of_Output = numOutputElements * sizeof(float);
70
71
      hostOutput = (float *)malloc(numOutputElements * sizeof(float));
72
     wbTime_stop(Generic, "Importing data and creating memory on host");
73
74
75
     wbLog(TRACE, "The number of input elements in the input is ", numInputE
     wbLog(TRACE, "The number of output elements in the input is ", numOutpu
76
77
     wbTime_start(GPU, "Allocating GPU memory.");
78
79
      //@@ Allocate GPU memory here
```

webgpu.com/attempt/176392 3/5

```
cudaMalloc(&deviceInput, Size_of_Input);
80
81
      cudaMalloc(&deviceOutput, Size_of_Output);
82
83
      wbTime_stop(GPU, "Allocating GPU memory.");
84
85
      wbTime_start(GPU, "Copying input memory to the GPU.");
      //@@ Copy memory to the GPU here
86
87
      cudaMemcpy(deviceInput, hostInput, Size_of_Input, cudaMemcpyHostToDevic
88
89
      wbTime_stop(GPU, "Copying input memory to the GPU.");
90
      //@@ Initialize the grid and block dimensions here
91
      dim3 dimGrid(numOutputElements, 1, 1);
92
      dim3 dimBlock(BLOCK_SIZE, 1, 1);
93
94
      wbTime_start(Compute, "Performing CUDA computation");
      //@@ Launch the GPU Kernel here
95
96
      reduction<<<dimGrid, dimBlock>>>(deviceInput, deviceOutput, numInputEle
97
98
      cudaDeviceSynchronize();
99
      wbTime_stop(Compute, "Performing CUDA computation");
100
101
      wbTime_start(Copy, "Copying output memory to the CPU");
102
      //@@ Copy the GPU memory back to the CPU here
103
      cudaMemcpy(hostOutput, deviceOutput, Size_of_Output, cudaMemcpyDeviceTo
104
105
      wbTime_stop(Copy, "Copying output memory to the CPU");
106
107
      /************************
108
       * Reduce output vector on the host
       * NOTE: One could also perform the reduction of the output vector
109
       * recursively and support any size input. For simplicity, we do not
110
       * require that for this lab.
111
       112
113
      for (ii = 1; ii < numOutputElements; ii++) {</pre>
114
        hostOutput[0] += hostOutput[ii];
115
      }
116
117
      wbTime_start(GPU, "Freeing GPU Memory");
      //@@ Free the GPU memory here
118
119
      cudaFree(deviceInput);
      cudaFree(deviceOutput);
120
121
122
      wbTime_stop(GPU, "Freeing GPU Memory");
123
      wbSolution(args, hostOutput, 1);
124
```

webgpu.com/attempt/176392 4/5

```
125

126  free(hostInput);

127  free(hostOutput);

128

129  return 0;

130 }
```

Designed and architected by Abdul Dakkak (https://www.dakkak.dev/).

webgpu.com/attempt/176392 5/5