

# List Reduction Grade

## Grade Summary (History) (/grade/history/5984)

Created:	<a href="#">less than a minute ago (2022-04-11 03:16:06 +0000 UTC)</a>
Total Score:	100 out of 100 points
Coding Score:	100 out of 100 points
Questions Score:	0

## Program Code

```
1 // MP Reduction
2 // Given a list (lst) of length n
3 // Output its sum = lst[0] + lst[1] + ... + lst[n-1];
4
5 #include <wb.h>
6
7 #define BLOCK\_SIZE 512 //@@ You can change this
8
9 #define wbCheck\(stmt\)
10     do {
11         cudaError_t err = stmt;
12         if (err != cudaSuccess) {
13             wbLog(ERROR, "Failed to run stmt ", #stmt);
14             wbLog(ERROR, "Got CUDA error ... ", cudaGetErrorString(err));
15             return -1;
16         }
17     } while (0)
18
19 \_\_global\_\_ void reduction(float *input, float *output, int len) {
20     //@@ Load a segment of the input vector into shared memory
21     \_\_shared\_\_ float sum[BLOCK_SIZE * 2];
```

```
22
23     int thread_x  = threadIdx.x;
24
25     int load_in = 2*(BLOCK_SIZE * blockIdx.x) + thread_x;
26
27     sum[thread_x] = load_in < len ? input[load_in] : 0.0;
28     if (load_in + BLOCK_SIZE < len)
29         sum[BLOCK_SIZE + thread_x] = input[load_in + BLOCK_SIZE];
30     else
31         sum[BLOCK_SIZE + thread_x] = 0.0;
32
33     ///@@ Traverse the reduction tree
34     for (int travel_tree = BLOCK_SIZE; travel_tree >= 1; travel_tree >>= 1)
35         __syncthreads();
36         if (thread_x < travel_tree)
37             sum[thread_x] += sum[thread_x + travel_tree];
38     }
39
40     ///@@ Write the computed sum of the block to the output vector at the
41     ///@@ correct index
42     if (thread_x == 0){
43         output[blockIdx.x] = sum[0];
44     }
45 }
46
47 int main(int argc, char **argv) {
48     int ii;
49     wbArg_t args;
50     float *hostInput; // The input 1D list
51     float *hostOutput; // The output list
52     float *deviceInput;
53     float *deviceOutput;
54     int numInputElements; // number of elements in the input list
55     int numOutputElements; // number of elements in the output list
56
57     args = wbArg_read(argc, argv);
58
59     wbTime_start(Generic, "Importing data and creating memory on host");
60     hostInput = (float *)wbImport(wbArg_getInputFile(args, 0), &numInputEle
61
62     numOutputElements = numInputElements / (BLOCK_SIZE << 1);
63     if (numInputElements % (BLOCK_SIZE << 1)){
64         numOutputElements++;
65     }
66
```

```
67 //2 variables to help with sizes
68 int Size_of_Input = numInputElements * sizeof(float);
69 int Size_of_Output = numOutputElements * sizeof(float);
70
71 hostOutput = (float *)malloc(numOutputElements * sizeof(float));
72
73 wbTime_stop(Generic, "Importing data and creating memory on host");
74
75 wbLog	TRACE, "The number of input elements in the input is ", numInputE
76 wbLog	TRACE, "The number of output elements in the input is ", numOutpu
77
78 wbTime_start(GPU, "Allocating GPU memory.");
79 //@@ Allocate GPU memory here
80 cudaMalloc(&deviceInput, Size_of_Input);
81 cudaMalloc(&deviceOutput, Size_of_Output);
82
83 wbTime_stop(GPU, "Allocating GPU memory.");
84
85 wbTime_start(GPU, "Copying input memory to the GPU.");
86 //@@ Copy memory to the GPU here
87 cudaMemcpy(deviceInput, hostInput, Size_of_Input, cudaMemcpyHostToDevice);
88
89 wbTime_stop(GPU, "Copying input memory to the GPU.");
90 //@@ Initialize the grid and block dimensions here
91 dim3 dimGrid(numOutputElements, 1, 1);
92 dim3 dimBlock(BLOCK_SIZE, 1, 1);
93
94 wbTime_start(Compute, "Performing CUDA computation");
95 //@@ Launch the GPU Kernel here
96 reduction<<<dimGrid, dimBlock>>>(deviceInput, deviceOutput, numInputEle
97
98 cudaDeviceSynchronize();
99 wbTime_stop(Compute, "Performing CUDA computation");
100
101 wbTime_start(Copy, "Copying output memory to the CPU");
102 //@@ Copy the GPU memory back to the CPU here
103 cudaMemcpy(hostOutput, deviceOutput, Size_of_Output, cudaMemcpyDeviceToHost);
104
105 wbTime_stop(Copy, "Copying output memory to the CPU");
106
107 /*****
108  * Reduce output vector on the host
109  * NOTE: One could also perform the reduction of the output vector
110  * recursively and support any size input. For simplicity, we do not
111  * require that for this lab.
```

```
112      *****/
113      for (ii = 1; ii < numOutputElements; ii++) {
114          hostOutput[0] += hostOutput[ii];
115      }
116
117      wbTime_start(GPU, "Freeing GPU Memory");
118      //@@ Free the GPU memory here
119      cudaFree(deviceInput);
120      cudaFree(deviceOutput);
121
122      wbTime_stop(GPU, "Freeing GPU Memory");
123
124      wbSolution(args, hostOutput, 1);
125
126      free(hostInput);
127      free(hostOutput);
128
129      return 0;
130 }
131
```