

◀ (/mp/5985?show=code) List Scan Attempt

Attempt Summary

Submit Attempt for Grading

Remember to answer the questions before clicking.

Dataset Id:	5
Created:	<u>less than a minute ago (2022-04-18 04:22:14 +0000 UTC)</u>
Status:	Correct solution for this dataset.

Timer Output

Kind	Location	Time (ms)	Message
Generic	main.cu::147	41.303532	Importing data and creating memory on host
GPU	main.cu::154	1.119895	Allocating GPU memory.
GPU	main.cu::161	0.028025	Clearing output memory.
GPU	main.cu::165	0.048714	Copying input memory to the GPU.
Compute	main.cu::173	0.072838	Performing CUDA computation
Copy	main.cu::187	0.039268	Copying output memory to the CPU
GPU	main.cu::191	0.167033	Freeing GPU Memory

Logger Output

Level	Location	Message
Trace	main::152	The number of input elements in the input is 9010

Program Code

```

1  /*
2  Dataset Id: 5
3  Created:  less than a minute ago
4  Status: Correct solution for this dataset.
5  Timer Output
6  Kind  Location  Time (ms)  Message
7  Generic main.cu::127  54.583396  Importing data and creating memory on ho
8  GPU main.cu::134  1.679192  Allocating GPU memory.
9  GPU main.cu::141  0.043779  Clearing output memory.
10 GPU main.cu::145  0.064188  Copying input memory to the GPU.
11 Compute main.cu::153  0.092417  Performing CUDA computation
12 Copy main.cu::167  0.05866  Copying output memory to the CPU
13 GPU main.cu::171  0.2507  Freeing GPU Memory
14 Logger Output
15 Level Location  Message
16 Trace main::132 The number of input elements in the input is 9010
17
18 */
19
20
21 // MP Scan
22 // Given a list (lst) of length n
23 // Output its prefix sum = {lst[0], lst[0] + lst[1], lst[0] + lst[1] + ..
24 // +
25 // lst[n-1]}
26
27 #include <wb.h>
28
29 #define BLOCK_SIZE 512 //@@ You can change this
30
31 #define wbCheck(stmt)
32 do {

```

```
33     cudaError_t err = stmt;
34     if (err != cudaSuccess) {
35         wbLog(ERROR, "Failed to run stmt ", #stmt);
36         wbLog(ERROR, "Got CUDA error ... ", cudaGetErrorString(err));
37         return -1;
38     }
39 } while (0)
40
41 __global__ void scan(float *input, float *output, int len, int flag) {
42     //@@ Modify the body of this function to complete the functionality of
43     //@@ the scan on the device
44     //@@ You may need multiple kernel calls; write your kernels before this
45     //@@ function and call them from here
46     __shared__ float block_array_scan[2 * BLOCK_SIZE];
47
48     int loop_index = 0;
49     int stride = 1;
50     if (!flag){
51         loop_index = (2 * blockIdx.x * blockDim.x) + threadIdx.x;
52         stride = blockDim.x;
53     }
54     else{
55         loop_index = (threadIdx.x + 1) * (2 * blockDim.x) - 1;
56         stride = 2 * blockDim.x;
57     }
58
59     int storeIndex = (2 * blockIdx.x * blockDim.x) + threadIdx.x;
60
61     //data input
62     if (loop_index < len){
63         block_array_scan[threadIdx.x] = input[loop_index];
64     }
65     else{
66         block_array_scan[threadIdx.x] = 0;
67     }
68     if (loop_index + stride < len){
69         block_array_scan[threadIdx.x + blockDim.x] = input[loop_index + stri
70     }
71     else{
72         block_array_scan[threadIdx.x + blockDim.x] = 0;
73     }
74
75     //First Step: Reduction
76     for (int stride = 1; stride <= (2 * BLOCK_SIZE); stride *= 2) {
77         __syncthreads();
```

```
78
79     int loop_index = (threadIdx.x + 1) * 2 * stride - 1;
80
81     if ((loop_index < 2 * BLOCK_SIZE) && ((loop_index - stride) >= 0)){
82         block_array_scan[loop_index] += block_array_scan[loop_index - s
83     }
84 }
85
86 //Use Distribution Tree method after Scanning
87 for (int stride = 2 * BLOCK_SIZE / 4; stride > 0; stride /= 2) {
88     __syncthreads();
89
90     int loop_index = (threadIdx.x + 1) * 2 * stride - 1;
91     if ((loop_index + stride) < 2 * BLOCK_SIZE){
92         block_array_scan[loop_index + stride] += block_array_scan[loop_
93     }
94 }
95
96 __syncthreads();
97 if (storeIndex < len){
98     output[storeIndex] = block_array_scan[threadIdx.x];
99 }
100
101 if (storeIndex + blockDim.x < len){
102     output[storeIndex + blockDim.x] = block_array_scan[threadIdx.x + blo
103 }
104 }
105
106 __global__ void add(float *input, float *output, float *array_sum, int le
107     __shared__ float move_loop;
108
109     int loop_index = threadIdx.x + (2 * blockIdx.x * blockDim.x);
110
111     if (threadIdx.x == 0){
112         if (blockIdx.x == 0){
113             move_loop = 0;
114         }
115         else{
116             move_loop = array_sum[blockIdx.x - 1];
117         }
118     }
119
120     __syncthreads();
121
122     if (loop_index < len){
```

```
123     output[loop_index] = input[loop_index] + move_loop;
124 }
125 if (loop_index + blockDim.x < len){
126     output[loop_index + blockDim.x] = input[loop_index + blockDim.x] + mo
127 }
128 }
129
130
131 int main(int argc, char **argv) {
132     wbArg_t args;
133     float *hostInput; // The input 1D list
134     float *hostOutput; // The output list
135     float *deviceInput;
136     float *deviceOutput;
137     int numElements; // number of elements in the list
138
139     //Additional Variables
140     //store temporary results from scanning
141     //store block summations from scanning
142     float *device_temporary_value;
143     float *scanned_dev_temp_val;
144
145     args = wbArg_read(argc, argv);
146
147     wbTime_start(Generic, "Importing data and creating memory on host");
148     hostInput = (float *)wbImport(wbArg_getInputFile(args, 0), &numElements
149     hostOutput = (float *)malloc(numElements * sizeof(float));
150     wbTime_stop(Generic, "Importing data and creating memory on host");
151
152     wbLog	TRACE, "The number of input elements in the input is ", numElemen
153
154     wbTime_start(GPU, "Allocating GPU memory.");
155     wbCheck(cudaMalloc((void **)&deviceInput, numElements * sizeof(float)))
156     wbCheck(cudaMalloc((void **)&deviceOutput, numElements * sizeof(float))
157     wbCheck(cudaMalloc((void **)&device_temporary_value, numElements * siz
158     wbCheck(cudaMalloc((void **)&scanned_dev_temp_val, 2 * BLOCK_SIZE * siz
159     wbTime_stop(GPU, "Allocating GPU memory.");
160
161     wbTime_start(GPU, "Clearing output memory.");
162     wbCheck(cudaMemset(deviceOutput, 0, numElements * sizeof(float)));
163     wbTime_stop(GPU, "Clearing output memory.");
164
165     wbTime_start(GPU, "Copying input memory to the GPU.");
166     wbCheck(cudaMemcpy(deviceInput, hostInput, numElements * sizeof(float),
167     wbTime_stop(GPU, "Copying input memory to the GPU.");
```

```
168
169 //@@ Initialize the grid and block dimensions here
170 dim3 dimGrid(ceil(numElements/(BLOCK_SIZE * 2.0)), 1, 1);
171 dim3 dimBlock(BLOCK_SIZE, 1, 1);
172
173 wbTime_start(Compute, "Performing CUDA computation");
174 //@@ Modify this to complete the functionality of the scan
175 //@@ on the device
176
177 //Here I store the temporary value in deviceOutput
178 scan<<<dimGrid, dimBlock>>>(deviceInput, device_temporary_value, numEle
179
180 dim3 postScanGrid(1, 1, 1);
181 scan<<<postScanGrid, dimBlock>>>(device_temporary_value, scanned_dev_te
182 add<<<dimGrid, dimBlock>>>(device_temporary_value, deviceOutput, scanne
183
184 cudaDeviceSynchronize();
185 wbTime_stop(Compute, "Performing CUDA computation");
186
187 wbTime_start(Copy, "Copying output memory to the CPU");
188 wbCheck(cudaMemcpy(hostOutput, deviceOutput, numElements * sizeof(float
189 wbTime_stop(Copy, "Copying output memory to the CPU");
190
191 wbTime_start(GPU, "Freeing GPU Memory");
192 cudaFree(deviceInput);
193 cudaFree(deviceOutput);
194 cudaFree(device_temporary_value);
195 cudaFree(scanned_dev_temp_val);
196 wbTime_stop(GPU, "Freeing GPU Memory");
197
198 wbSolution(args, hostOutput, numElements);
199
200 free(hostInput);
201 free(hostOutput);
202
203 return 0;
204 }
```

