

## ◀ (/mp/124?show=code) Sparse Matrix Multiplication (JDS) Attempt

### Attempt Summary

Submit Attempt for Grading

Remember to answer the questions before clicking.

Dataset Id:	6
Created:	<u>about a minute ago (2022-04-24 19:03:15 +0000 UTC)</u>
Status:	Correct solution for this dataset.

### Timer Output

Kind	Location	Time (ms)	Message
<b>Generic</b>	main.cu::65	489.387291	Importing data and creating memory on host
<b>GPU</b>	main.cu::78	1.178791	Allocating GPU memory.
<b>GPU</b>	main.cu::89	0.25224	Copying input memory to the GPU.
<b>Compute</b>	main.cu::98	0.170311	Performing CUDA computation
<b>Copy</b>	main.cu::103	0.031733	Copying output memory to the CPU
<b>GPU</b>	main.cu::107	0.23004	Freeing GPU Memory

## Program Code

```
1  #include <wb.h>
2
3  #define wbCheck(stmt)
4      do {
5          cudaError_t err = stmt;
6          if (err != cudaSuccess) {
7              wbLog(ERROR, "Failed to run stmt ", #stmt);
8              wbLog(ERROR, "Got CUDA error ... ", cudaGetErrorString(err));
9              return -1;
10         }
11     } while (0)
12
13 __global__ void spmvJDSKernel(float *out, int *matColStart, int *matCols,
14     //@ insert spmv kernel for jds format
15
16     //1D array method
17     int row = blockIdx.x * blockDim.x + threadIdx.x;
18     if (row < dim){
19         float prod = 0;
20         for (int i = 0; i < matRows[row]; i++){
21             prod += matData[matColStart[i] + row] * vec[matCols[matColStart[i]
22         }
23
24         //output
25         out[matRowPerm[row]] = prod;
26     }
27 }
28
29 static void spmvJDS(float *out, int *matColStart, int *matCols, int *matR
30
31     //@ invoke spmv kernel for jds format
32
33     //dimensions
34     dim3 grid(ceil((float)dim/512.0), 1, 1);
35     dim3 block(512, 1, 1);
36
37     //call kernel
38     spmvJDSKernel<<<grid, block>>>(out, matColStart, matCols, matRowPerm, m
39 }
40
41 int main(int argc, char **argv) {
```

```
42  wbArg_t args;
43  int *hostCSRCols;
44  int *hostCSRRows;
45  float *hostCSRData;
46  int *hostJDSColStart;
47  int *hostJDSCols;
48  int *hostJDSRowPerm;
49  int *hostJDSRows;
50  float *hostJDSData;
51  float *hostVector;
52  float *hostOutput;
53  int *deviceJDSColStart;
54  int *deviceJDSCols;
55  int *deviceJDSRowPerm;
56  int *deviceJDSRows;
57  float *deviceJDSData;
58  float *deviceVector;
59  float *deviceOutput;
60  int dim, ncols, nrows, ndata;
61  int maxRowNNZ;
62
63  args = wbArg_read(argc, argv);
64
65  wbTime_start(Generic, "Importing data and creating memory on host");
66  hostCSRCols = (int *)wbImport(wbArg_getInputFile(args, 0), &ncols, "Int
67  hostCSRRows = (int *)wbImport(wbArg_getInputFile(args, 1), &nrows, "Int
68  hostCSRData = (float *)wbImport(wbArg_getInputFile(args, 2), &ndata, "R
69  hostVector = (float *)wbImport(wbArg_getInputFile(args, 3), &dim, "Real
70
71  hostOutput = (float *)malloc(sizeof(float) * dim);
72
73  wbTime_stop(Generic, "Importing data and creating memory on host");
74
75  CSRTtoJDS(dim, hostCSRRows, hostCSRCols, hostCSRData, &hostJDSRowPerm, &
76  maxRowNNZ = hostJDSRows[0];
77
78  wbTime_start(GPU, "Allocating GPU memory.");
79  cudaMalloc((void **)&deviceJDSColStart, sizeof(int) * maxRowNNZ);
80  cudaMalloc((void **)&deviceJDSCols, sizeof(int) * ndata);
81  cudaMalloc((void **)&deviceJDSRowPerm, sizeof(int) * dim);
82  cudaMalloc((void **)&deviceJDSRows, sizeof(int) * dim);
83  cudaMalloc((void **)&deviceJDSData, sizeof(float) * ndata);
84
85  cudaMalloc((void **)&deviceVector, sizeof(float) * dim);
86  cudaMalloc((void **)&deviceOutput, sizeof(float) * dim);
```

```
87  wbTime_stop(GPU, "Allocating GPU memory.");
88
89  wbTime_start(GPU, "Copying input memory to the GPU.");
90  cudaMemcpy(deviceJDSColStart, hostJDSColStart, sizeof(int) * maxRowNNZ,
91  cudaMemcpy(deviceJDSCols, hostJDSCols, sizeof(int) * ndata, cudaMemcpyH
92  cudaMemcpy(deviceJDSRowPerm, hostJDSRowPerm, sizeof(int) * dim, cudaMemcpy
93  cudaMemcpy(deviceJDSRows, hostJDSRows, sizeof(int) * dim, cudaMemcpyHos
94  cudaMemcpy(deviceJDSData, hostJDSData, sizeof(float) * ndata, cudaMemcpy
95  cudaMemcpy(deviceVector, hostVector, sizeof(float) * dim, cudaMemcpyHos
96  wbTime_stop(GPU, "Copying input memory to the GPU.");
97
98  wbTime_start(Compute, "Performing CUDA computation");
99  spmvJDS(deviceOutput, deviceJDSColStart, deviceJDSCols, deviceJDSRowPer
100 cudaMemcpyDeviceSynchronize();
101 wbTime_stop(Compute, "Performing CUDA computation");
102
103 wbTime_start(Copy, "Copying output memory to the CPU");
104 cudaMemcpy(hostOutput, deviceOutput, sizeof(float) * dim, cudaMemcpyDev
105 wbTime_stop(Copy, "Copying output memory to the CPU");
106
107 wbTime_start(GPU, "Freeing GPU Memory");
108 cudaFree(deviceVector);
109 cudaFree(deviceOutput);
110 cudaFree(deviceJDSColStart);
111 cudaFree(deviceJDSCols);
112 cudaFree(deviceJDSRowPerm);
113 cudaFree(deviceJDSRows);
114 cudaFree(deviceJDSData);
115
116 wbTime_stop(GPU, "Freeing GPU Memory");
117
118 wbSolution(args, hostOutput, dim);
119
120 free(hostCSRCols);
121 free(hostCSRRows);
122 free(hostCSRData);
123 free(hostVector);
124 free(hostOutput);
125 free(hostJDSColStart);
126 free(hostJDSCols);
127 free(hostJDSRowPerm);
128 free(hostJDSRows);
129 free(hostJDSData);
130
131 return 0;
```

132  
133

}